

Ethics Worksheet

Conor Houghton

1 Ethics

The goal in this workshop is to think about the ethics of AI and machine learning. Please divide into groups of three or four and discuss these topics; ideally write a short response to each.

1. **Who owns a large machine learning system?** Does a large machine learning system belong to the company that developed it? We all created the data used to train models and we all own the historic scientific advances the models exploit; do the models belong to us all?
2. **Who bears responsibility when AI systems cause harm?** Should accountability fall on the developers who built the system, the companies that deployed it, the users who operated it, or the policymakers who failed to regulate it. Does the answer change depending on how autonomous the system is?
3. **Can an AI system ever be truly “fair”?** Since AI models are trained on human-generated data that reflects historical inequalities, is it possible to build systems that don’t perpetuate or amplify existing biases, or is some degree of encoded bias unavoidable?
4. **Where should the line be drawn between beneficial surveillance and privacy violation?** AI-powered tools can monitor workers, predict criminal behaviour, or track health outcomes in ways that may benefit society. At what point do these applications cross an ethical boundary?
5. **Should there be domains where AI decision-making is categorically off-limits?** Areas like criminal sentencing, medical diagnosis, or military targeting involve high-stakes judgments about human lives. Are there decisions that should remain with humans regardless of how accurate or efficient an AI system becomes?
6. **Who gets to define “aligned” AI, and aligned with whose values?** The goal of building AI that reflects human values assumes a degree of consensus that may not exist across cultures, political systems, and socioeconomic contexts. How should the AI development community navigate the question of whose ethics get encoded into these systems?
7. **Do we have, or will we ever have, moral obligations to AI systems?** If a sufficiently advanced AI were to exhibit behaviour indistinguishable from consciousness, suffering, or preference, would we be ethically obligated to extend it some form of rights or protections? Is this a question that science can settle or is it ultimately a matter of interpretation? Who has the authority to decide?